

MODEL SELECTION & TOKEN ECONOMICS

Haiku Claude 3.5 · 200K · \$Lowest

Summaries, classification, extraction, high-volume Q&A. Fastest response. Use for anything you'd send to an intern.

Lite Nova Lite · 128K · \$Very Low

Reformatting, basic lookups, repetitive batch tasks. Cheapest option when quality tolerance is high.

Sonnet Claude 4 · 200K · \$Moderate ★ START HERE

Complex analysis, policy interpretation, detailed writing, code review. Default for all knowledge work. Start here.

Pro Nova Pro · 128K · \$Moderate

Multimodal (text + images/video). Good for decks with visuals or image-heavy documents.

Opus Claude 4 · 200K · \$Highest

Regulatory interpretation, multi-step reasoning, agentic workflows, anything where a single error is costly.

TOKEN CHEAT SHEET

1 token	≈ ¾ of a word
200-word email	≈ 270 tokens
10-page report	≈ 6,700 tokens
RAG injection (5 chunks)	≈ 2,500 tokens
Max output / response	≈ 48,000 words
Switch models mid-chat	@ModelName in chat

KNOWLEDGE BASE MASTERY

How RAG works — what actually happens when you upload a file:

1. Upload

Apache Tika extracts text from 1,000+ formats (PDF, Word, Excel, PPT, HTML, email, ZIP...)

2. Chunk

Document split into ~1,500-character segments with configurable overlap to preserve context across boundaries

3. Embed

Each chunk converted to a vector (Cohere Embed v4) — a mathematical fingerprint of its meaning

4. Store

Vectors saved in PGVector (PostgreSQL) inside Flagstar's AWS — fully encrypted, never external

5. Retrieve

Your question is also embedded; system finds top 5 most semantically similar chunks

6. Answer

Those 5 chunks injected into the AI's context window with citations → grounded, cited response

Retrieval Modes

Focused Retrieval (default)

Sends only the top 5 most relevant chunks (~2,500 tokens). Best for large document collections. Efficient and fast.

ADVANCED PROMPT ARCHITECTURE

The CRAFT Framework

C Context

Background, constraints, who you are, what this is for

R Role

"You are a senior Treasury analyst at Flagstar..." — primes the model's perspective and expertise level

A Action

Precise verb: analyze / draft / compare / extract / evaluate. Avoid: 'help me with' or 'tell me about'

F Format

"Return as a markdown table with columns X,Y,Z" or "3 bullet points, max 25 words each" — be exact

T Tone & Target

Audience, reading level, formality. "Plain English for a non-finance VP" changes everything

POWER TECHNIQUES

Chain-of-thought

Add "think step by step" or "show your reasoning" to any analytical prompt. Dramatically improves accuracy on multi-step problems.

ANTI-PATTERNS

- ✗ Vague ask: "Help me with this" → be explicit about task, format, and audience
- ✗ Using Opus for simple summaries → wasteful; switch to Haiku
- ✗ Ignoring the first response → iterate, don't restart from scratch
- ✗ No context → AI has no idea who will read this or why
- ✗ Uploading massive files with no specific question → scopes nothing
- ✗ Treating output as final → always review before any official use
- ✗ Long context

KNOWLEDGE BASE BEST PRACTICES

Organize by domain

Create separate KBs: "Compliance Policies", "ALM Procedures", "ALCO Minutes". Focused scope = higher retrieval precision.

Name & describe precisely

The description trains the AI's discovery. Write it like a search query: "Flagstar FTP methodology, ALM policies, interest rate risk procedures updated 2024–2025."

Keep documents current

Outdated documents produce outdated answers. Assign a KB owner per department and run a quarterly review cycle.

Test before sharing

Ask 5–10 representative questions and verify citations pull from the right chunks. Adjust document structure if retrieval is weak.

Scope your agents

When building a custom model, attach only the KBs it needs. Unscoped models can retrieve from any KB — increases hallucination risk.

Prefer text-heavy files

Charts, images, and complex table formatting are NOT extracted. Excel formulas show only computed values. Scanned PDFs require OCR.

Use overlap on flowing docs

For documents where meaning spans section boundaries (legal contracts, multi-part analyses), chunk overlap preserves context continuity.

BUILDING CUSTOM AI AGENTS

A Custom Model in Workspace is not a new AI — it is a configuration preset that binds a system prompt, knowledge bases, tools, skills, and capability toggles to a base model. Think of it as a purpose-built team member with a defined job, scope, and personality.

Example: Treasury Policy Assistant

Base Model:

StarIQ Sonnet — strong reasoning, cost-efficient for daily queries

System Prompt:

"You are a Treasury policy expert at Flagstar. Answer questions about FTP, ALM, and interest rate risk using ONLY attached knowledge bases. Always cite your sources. Say 'I don't know' rather than guessing. Today: {{CURRENT_DATE}}, user: {{USER_NAME}}."

Knowledge:

Attach: 'FTP Procedures', 'ALM Policy', 'ALCO Minutes' KBs

Capabilities:

Enable: Citations. Disable: Image Generation, Code Interpreter

Access:

Share with Treasury RBAC group — Read permission

Prompt chips:

Add starter suggestions: "Summarize FTP methodology", "What is

DOCUMENT GENERATION & IN-CHAT RENDERING

Generate downloadable Office files directly in chat:

SKILLS vs TOOLS vs FUNCTIONS

Word (.docx) — Reports, summaries, proposals, meeting notes

Skills

PowerPoint (.pptx)

Markdown instructions — Presentations with speaker notes, layouts, themes

Excel (.xlsx)

Work-on-Doc — Data exports, financial models, tracking sheets

In-Chat Artifacts (no download needed):

Python scripts — give the model access to real-world data, APIs, and calculations. Require security code review before deploy.

Mermaid

Flowcharts, org charts, Gantt, sequence, ERD, pie charts

HTML

Interactive dashboards, calculators, styled reports

Functions

SVG — Infographics, logos, technical drawings

Admin-only — add new AI providers, UI buttons, data filters. Not available to end users.

CSV/JSON

Tables and structured data with syntax highlighting

LaTeX

Mathematical equations and scientific notation

MULTI-MODEL WORKFLOW (example)

> @Haiku — extract all dates, amounts, party names from contract

> @Sonnet — analyze extracted terms for risk and ambiguity

> @Opus — draft legal memo with full reasoning and

2026 ROADMAP

Documents available for download for 30 days after generation.

Core platform, all models, KBs & RAG, prompt libraries, voice, accessibility

Q2 Apr–Jun 2026

OneDrive & SharePoint RAG integration · Mobile PWA · AI Web Search

(Palo Alto Prisma)

Q3 Jul–Sep 2026

Model Context Protocol · Python AI Tool Framework · Code Interpreter

enhancements

Q4 Oct–Dec 2026

RBAC & SHARING STRATEGY

All Workspace resources (Models, Knowledge, Prompts, Skills, Tools) share a unified, additive permission system. Permissions NEVER deny — if any group grants access, the user has it.

Two Group Types — Use Both

Permission Groups

Grant feature rights (image gen, code interpreter, web search).

Set 'Who can share' = Nobody. Example: 'Power Users',

'Developers'.

Sharing Groups

Organize teams for resource access. Set 'Who can share' =

Members or Anyone. Example: 'Treasury Team', 'HR Leadership',

'Risk Management'.

PRINCIPLE OF LEAST PRIVILEGE

- Minimize Global Default Permissions — grant rights via specific groups, not globally
- Attach only the KBs an agent needs — unscoped models can retrieve from all user KBs
- Use Read access for team sharing; Write access only for designated owners
- KBs with confidential content → Private visibility + explicit access grants only

DATA GOVERNANCE & COMPLIANCE

Roles: Admin → full access | User → standard | Pending → New SSO accounts
Annual reports, regulatory guidance, press releases — freely uploadable

INTERNAL LIBRARY WITH SLASH COMMANDS

Create reusable templates with project-level variables for type

/ in StarIQ to trigger your template — a form appears, they fill it in,

output is structured.

Confidential

Strategy docs, non-public financials — configure KB access

controls first

{{project_name}} select:...}} Dropdown: High / Medium / Low

{{deadline | date:required}} Date picker, required

{{urgent | checkbox}} Boolean toggle

Customer PII, SSNs, account numbers, credentials — DO NOT

upload without explicit authorization

{{CURRENT_DATE}} Auto-injected — today's date

{{USER_NAME}} Auto-injected — requestor's name

USE THE AI INTAKE PROCESS WHEN:

- Building agents that access sensitive or regulated data
- Integrating StarIQ with other Flagstar systems or databases
- Deploying AI for customer-facing processes

TROUBLESHOOTING QUICK REFERENCE

- Any initiative involving GLBA, SOX, HMDA regulated data
- SSO login fails
- Access via myapps.microsoft.com; check OAuth redirect URIs
- Using or reusing AI to process or mark

Everyday use (summarizing, drafting, analyzing) does NOT require intake.

Doc not found in KB

Check Documents tab for processing status; verify # reference

Outdated AI answers

Upload current docs; use web search when live (Q2 2026)